

Telecom Churn Prediction and Analysis

Objective:

The goal of this notebook is to explore and analyze the Telecom Churn dataset to understand factors contributing to customer churn and to develop a predictive model that can forecast customer churn with high accuracy.

```
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from sklearn.preprocessing import OneHotEncoder, StandardScaler
import pandas as pd
import numpy as np
```

Dataset Description:

The Telecom Churn dataset comprises customer data from a telecom company. Key features include customer account information, demographic data, service usage, and churn status (Yes or No).

Data Fields:

- CustomerID: Unique identifier for the customer
- Gender: Customer gender (Male, Female)
- Age: Customer age
- Tenure: Number of months the customer has stayed with the company
- ServiceCalls: Number of customer service calls made
- MonthlyCharges: The amount charged to the customer monthly
- TotalCharges: The total amount charged to the customer
- Churn: Customer churn status (Yes or No)

```
data = pd.read_csv("/content/Telecom_Customer_Churn_Dataset.csv")
data.head(10)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	M
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	
4	9237-HQITU	Female	0	No	No	2	Yes	
5	9305-CDSKC	Female	0	No	No	8	Yes	
6	1452-KIOVK	Male	0	No	Yes	22	Yes	
7	6713-OKOMC	Female	0	No	No	10	No	
8	7892-POOKP	Female	0	Yes	No	28	Yes	
9	6388-TABGU	Male	0	No	Yes	62	Yes	

10 rows × 21 columns

Data Preprocessing

```
data.shape

(7043, 21)

data.info()
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7032 non-null   float64
20  Churn                 7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

```
print(data['SeniorCitizen'].nunique())
data['SeniorCitizen'].unique()
```

```
→ 2
array([0, 1])
```

```
print(data['Contract'].nunique())
data['Contract'].unique()
```

```
→ 3
array(['Month-to-month', 'One year', 'Two year'], dtype=object)
```

```
print(data['PaymentMethod'].nunique())
data['PaymentMethod'].unique()
```

```
→ 4
array(['Electronic check', 'Mailed check', 'Bank transfer (automatic)',
      'Credit card (automatic)'], dtype=object)
```

```
print(data['Partner'].nunique())
data['Partner'].unique()
```

```
→ 2
array(['Yes', 'No'], dtype=object)
```

```
print(data['DeviceProtection'].nunique())
data['DeviceProtection'].unique()
```

```
→ 3
array(['No', 'Yes', 'No internet service'], dtype=object)
```

```
print(data['PhoneService'].nunique())
data['PhoneService'].unique()
```

```
→ 2
array(['No', 'Yes'], dtype=object)
```

```
print(data['Churn'].nunique())
data['Churn'].unique()
```

```
→ 2
array(['No', 'Yes'], dtype=object)
```

```
data.columns
```

```
→ Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```


```
data.isna().sum()
```

```
→ customerID      0
   gender         0
   SeniorCitizen  0
   Partner        0
```



```
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    11
Churn           0
dtype: int64
```

There are 11 null values in `TotalCharges` column which can be replaced with mean or median of the data column.

```
data.describe()
```




	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	0.162147	32.371149	64.761692	2283.300441
std	0.368612	24.559481	30.090047	2266.771362
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3794.737500
max	1.000000	72.000000	118.750000	8684.800000



```
# Calculate the mean of the column
mid = data['TotalCharges'].median()

# Replace missing values with the mean
data['TotalCharges'].fillna(mid, inplace=True)
```

```
data.isna().sum()
```



```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

```
data.duplicated().sum()
```



```
0
```

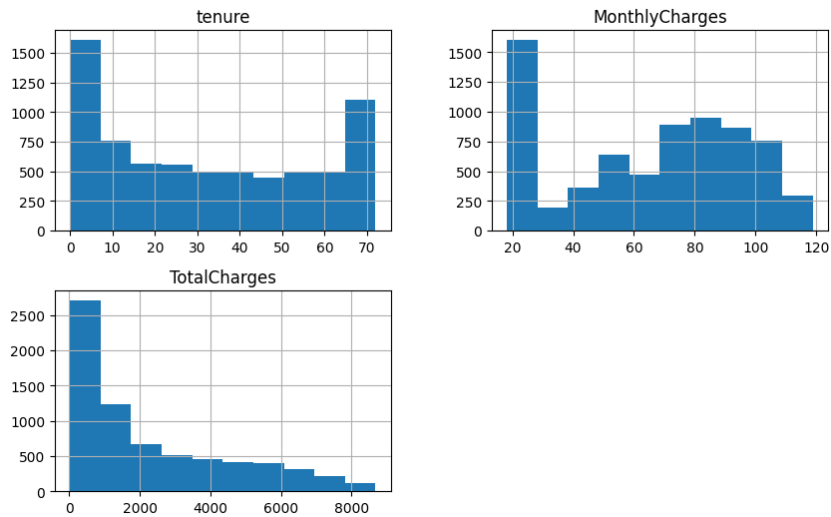
The dataset has no duplicates.

Exploratory Data Analysis

```
# Distribution of key variables
# Numerical variables
num_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
data[num_cols].hist(figsize=(10, 6))
plt.suptitle('Distribution of Numerical Variables')
plt.show()
```



Distribution of Numerical Variables




Categorical variables

```
cat_cols = ['gender', 'SeniorCitizen', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',  
            'InternetService', 'OnlineSecurity', 'DeviceProtection', 'TechSupport', 'StreamingTV',  
            'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'Churn']
```


```
for col in cat_cols:  
    plt.figure(figsize=(8, 4))  
    sns.countplot(data=data, x=col, hue='Churn')  
    plt.title(f'Distribution of {col} by Churn')  
    plt.xlabel(col)  
    plt.ylabel('Count')  
    plt.xticks(rotation=45)  
    plt.legend(title='Churn', loc='upper right')  
    plt.show()
```




Distribution of gender by Churn




Distribution of SeniorCitizen by Churn




Distribution of Partner by Churn

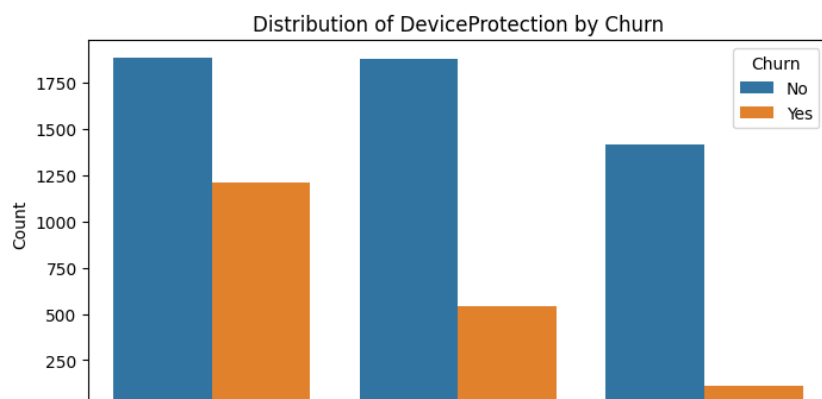


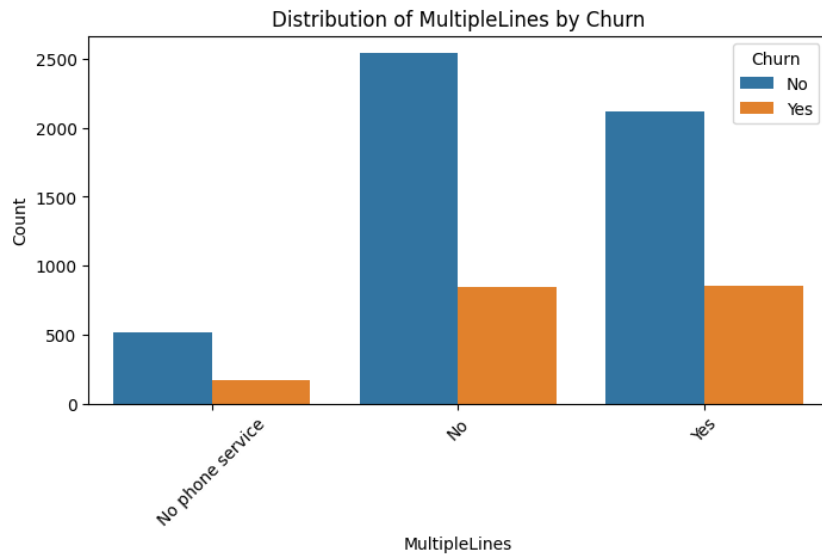




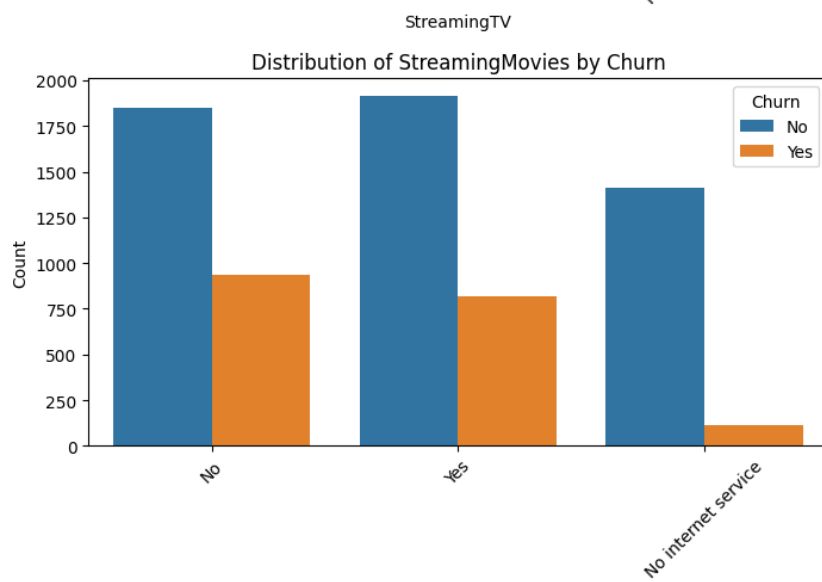
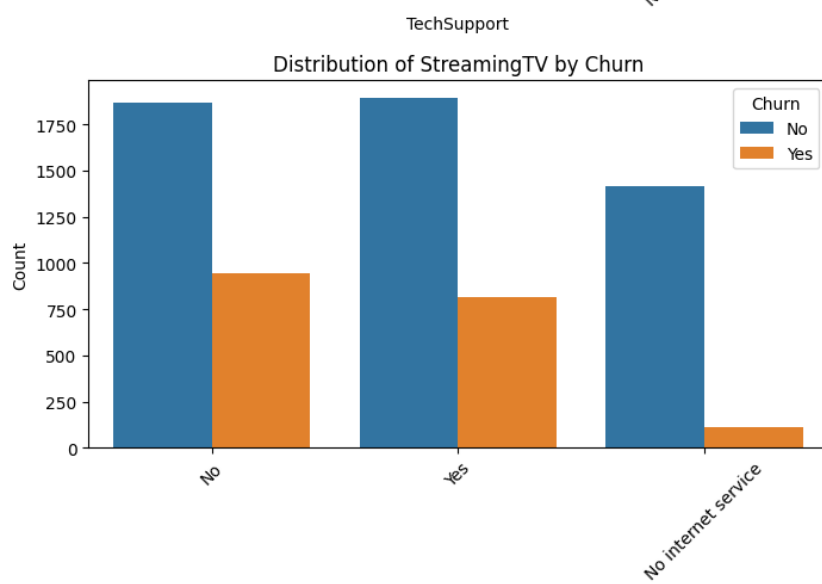
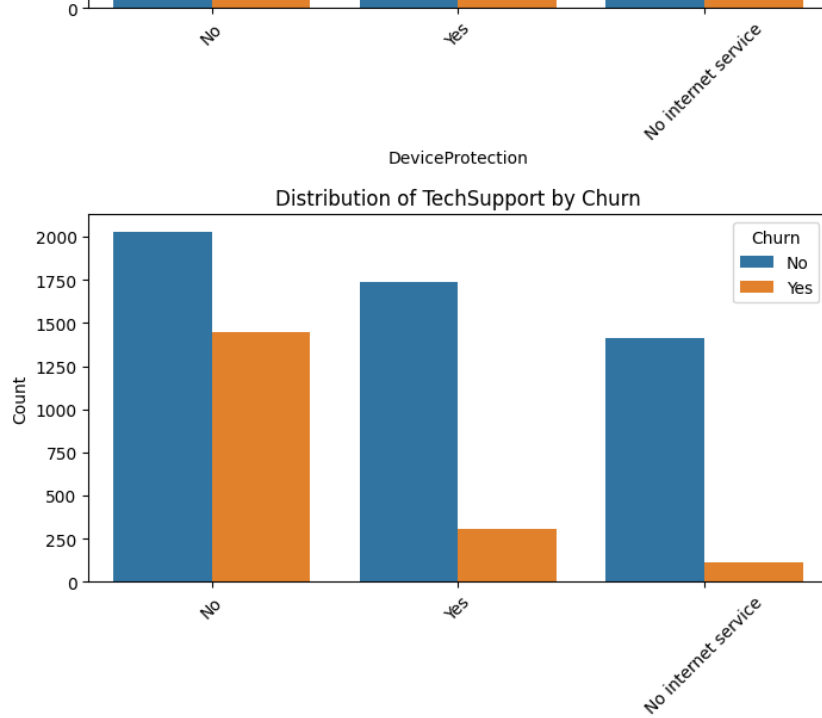
Distribution of Dependents by Churn


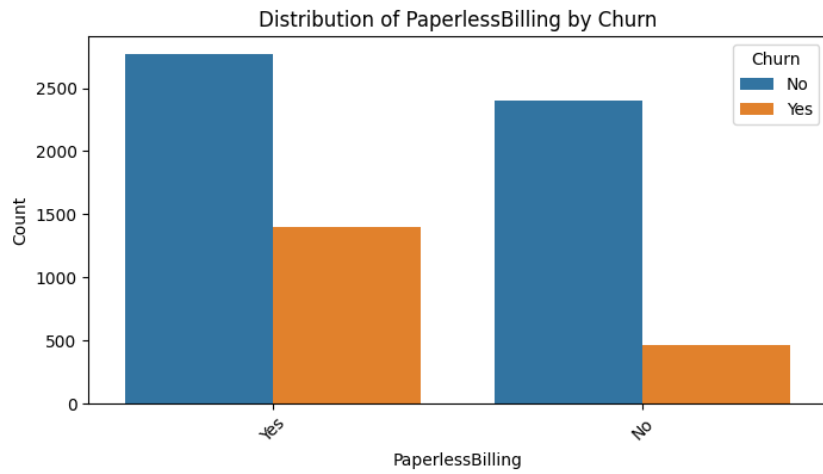



Distribution of PhoneService by Churn










WARNING:matplotlib.legend.No artists with labels found to put in legend. Note tha





```
# Investigate relationships between features
# Select only numeric columns
numeric_df = data.select_dtypes(include=['number'])

# Calculate correlation matrix
correlation_matrix = numeric_df.corr()
# Print the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)
# Correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```

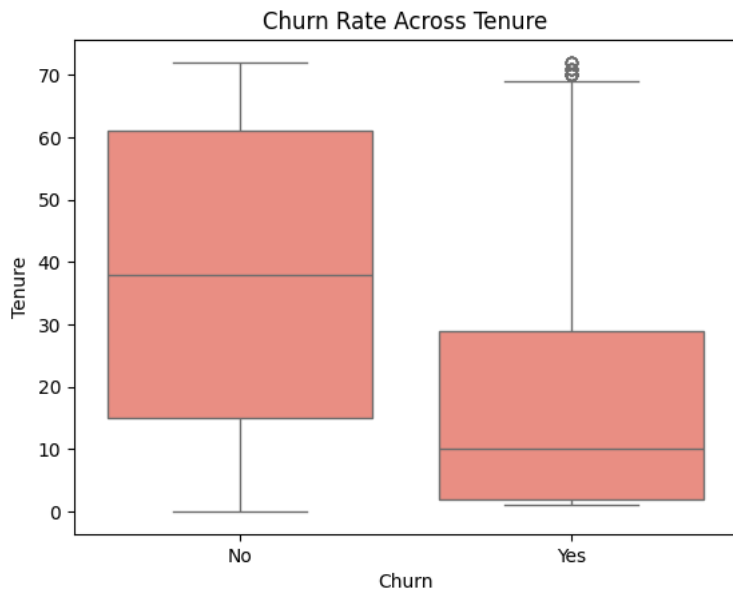
Correlation Matrix:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
SeniorCitizen	1.000000	0.016567	0.220173	0.102652
tenure	0.016567	1.000000	0.247900	0.825464
MonthlyCharges	0.220173	0.247900	1.000000	0.650864
TotalCharges	0.102652	0.825464	0.650864	1.000000



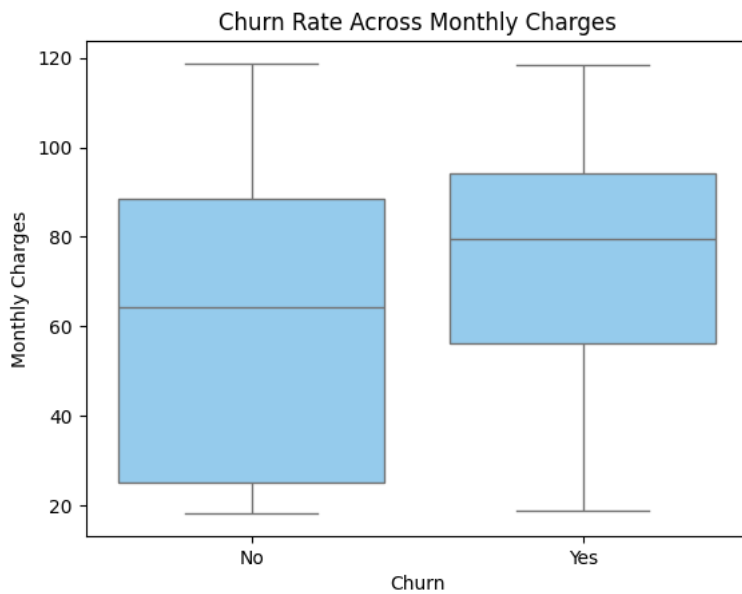
```
# Churn rate across tenure
sns.boxplot(x='Churn', y='tenure', data=data, color='salmon')
plt.title('Churn Rate Across Tenure')
plt.xlabel('Churn')
```

```
plt.ylabel('Tenure')
plt.show()
```




Long-term clients tend to stay while new clients having a short tenure of 1-2 years tend to churn.

```
# Churn rate across MonthlyCharges
sns.boxplot(x='Churn', y='MonthlyCharges', data=data, color='lightskyblue')
plt.title('Churn Rate Across Monthly Charges')
plt.xlabel('Churn')
plt.ylabel('Monthly Charges')
plt.show()
```



Increase in monthly charges is also seen to be a factor for higher churn rates.

```
# Churn rate across TotalCharges
sns.boxplot(x='Churn', y='TotalCharges', data=data, color='coral')
plt.title('Churn Rate Across Total Charges')
plt.xlabel('Churn')
plt.ylabel('Total Charges')
plt.show()
```



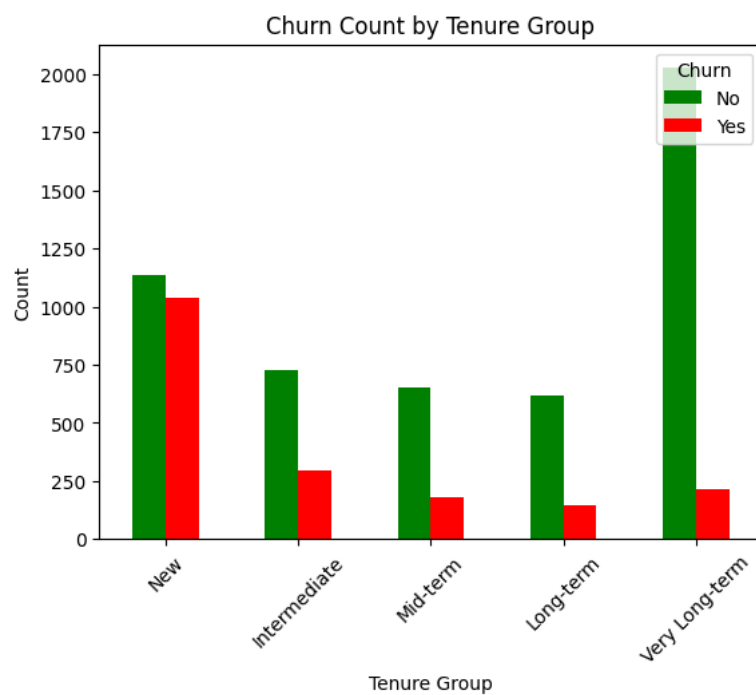
Total charges does not seem to be a significant factor for higher churn rate.

Feature Engineering

```
# Create new features
data['TenureGroup'] = pd.cut(data['tenure'], bins=[0, 12, 24, 36, 48, data['tenure'].max()], labels=['New', 'Intermediate', 'Mid-term', 'Long-term',
data['FamilyStatus'] = (data['Partner'] == 'Yes') | (data['Dependents'] == 'Yes')
data['NumServices'] = data[['PhoneService', 'InternetService', 'OnlineSecurity', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies']
```

```
# Calculate churn counts within each tenure group
grouped = data.groupby(['TenureGroup', 'Churn']).size().unstack()
```

```
# Plot grouped bar plot
grouped.plot(kind='bar', stacked=False, color=['green', 'red'])
plt.title('Churn Count by Tenure Group')
plt.xlabel('Tenure Group')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Churn', loc='upper right')
plt.show()
```



```
data.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	M
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes
2	3668-QPYBK	Male		0	No	No	2	Yes
3	7795-CFOCW	Male		0	No	No	45	No
4	9237-HQITU	Female		0	No	No	2	Yes

5 rows × 24 columns

```
# Interaction features
data['InternetAndPhoneService'] = (data['PhoneService'] == 'Yes') & (data['InternetService'] != 'No')
data['SecurityAndSupportBundle'] = (data['OnlineSecurity'] == 'Yes') & (data['TechSupport'] == 'Yes')
```

```
data.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	M
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes
2	3668-QPYBK	Male		0	No	No	2	Yes
3	7795-CFOCW	Male		0	No	No	45	No
4	9237-HQITU	Female		0	No	No	2	Yes

5 rows × 26 columns

Feature Encoding

```
# Encode categorical features
data = pd.get_dummies(data, columns=['Contract', 'PaymentMethod', 'MultipleLines', 'InternetService', 'PaymentMethod', 'TenureGroup'])
```

```
data.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	0
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes
2	3668-QPYBK	Male		0	No	No	2	Yes
3	7795-CFOCW	Male		0	No	No	45	No
4	9237-HQITU	Female		0	No	No	2	Yes

5 rows × 43 columns

```
from sklearn.preprocessing import LabelEncoder
# Select categorical columns
categorical_columns = ['gender', 'Partner', 'Dependents', 'PhoneService', 'OnlineSecurity', 'PaperlessBilling', 'DeviceProtection', 'OnlineBackup', 'Te

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Apply LabelEncoder to each categorical column
for col in categorical_columns:
    data[col] = label_encoder.fit_transform(data[col])
```

```
data.columns
```

```
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'OnlineSecurity', 'OnlineBackup',
      'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',
      'PaperlessBilling', 'MonthlyCharges', 'TotalCharges', 'Churn',
      'FamilyStatus', 'NumServices', 'InternetAndPhoneService',
      'SecurityAndSupportBundle', 'Contract_Month-to-month',
      'Contract_One year', 'Contract_Two year',
      'PaymentMethod_Bank transfer (automatic)',
      'PaymentMethod_Credit card (automatic)',
      'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check',
      'MultipleLines_No', 'MultipleLines_No phone service',
      'MultipleLines_Yes', 'InternetService_DSL',
      'InternetService_Fiber optic', 'InternetService_No',
      'PaymentMethod_Bank transfer (automatic)',
      'PaymentMethod_Credit card (automatic)',
      'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check',
      'TenureGroup_New', 'TenureGroup_Intermediate', 'TenureGroup_Mid-term',
      'TenureGroup_Long-term', 'TenureGroup_Very Long-term'],
      dtype='object')
```

```
data.drop(columns=['customerID'], inplace=True) # is not required for analysis
```


```
data['Churn'] = data['Churn'].replace({'Yes': 1, 'No': 0})
```

```
data.head()
```



	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	OnlineSecurit
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	

5 rows × 42 columns



```
# Feature scaling
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
```


```
data[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(data[['tenure', 'MonthlyCharges', 'TotalCharges']])
```

```
data.head()
```



	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	OnlineSecur:
0	0	0	1	0	0.013889	0	
1	1	0	0	0	0.472222	1	
2	1	0	0	0	0.027778	1	
3	1	0	0	0	0.625000	0	
4	0	0	0	0	0.027778	1	

5 rows × 42 columns



Feature Selection

```
# Feature selection
```

```
# Example: Recursive Feature Elimination (RFE) with a Random Forest Classifier
```

```
from sklearn.feature_selection import RFE
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
X = data.drop(columns=['Churn'])
```

```
y = data['Churn']
```

```
rf_model = RandomForestClassifier()
```

```
rfe = RFE(rf_model, n_features_to_select=10) # Select top 10 features
```

```
rfe.fit(X, y)
```

```
selected_features = X.columns[rfe.support_]
```

```
print("Selected Features:", selected_features)
```

```
➔ Selected Features: Index(['gender', 'tenure', 'OnlineSecurity', 'TechSupport', 'MonthlyCharges',
'TotalCharges', 'NumServices', 'Contract_Month-to-month',
'InternetService_Fiber optic', 'PaymentMethod_Electronic check'],
dtype='object')
```

We can reduce the dimensionality of the data by selecting the most important features. This will result in an optimum network and improved results.

▼ Predictive Modeling

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
```

```
# Logistic Regression
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)
```

```
➔ ▼ RandomForestClassifier
RandomForestClassifier()
```

```
# Random Forest Classifier
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
```

```
➔ ▼ RandomForestClassifier
RandomForestClassifier()
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
```

```
# Predictions
logistic_preds = logistic_model.predict(X_test)
rf_preds = rf_model.predict(X_test)
```

```
# Evaluation metrics
print("Logistic Regression:")
print("Accuracy:", accuracy_score(y_test, logistic_preds))
print("Precision:", precision_score(y_test, logistic_preds))
print("Recall:", recall_score(y_test, logistic_preds))
print("F1 Score:", f1_score(y_test, logistic_preds))
print("ROC-AUC Score:", roc_auc_score(y_test, logistic_preds))
```

```
print("\nRandom Forest Classifier:")
print("Accuracy:", accuracy_score(y_test, rf_preds))
print("Precision:", precision_score(y_test, rf_preds))
print("Recall:", recall_score(y_test, rf_preds))
print("F1 Score:", f1_score(y_test, rf_preds))
print("ROC-AUC Score:", roc_auc_score(y_test, rf_preds))
```

```
➔ Logistic Regression:
Accuracy: 0.8126330731014905
Precision: 0.6763754045307443
Recall: 0.5603217158176944
F1 Score: 0.6129032258064516
ROC-AUC Score: 0.731898309646299
```

```
Random Forest Classifier:
Accuracy: 0.7899219304471257
Precision: 0.6452830188679245
Recall: 0.4584450402144772
F1 Score: 0.5360501567398118
ROC-AUC Score: 0.6838557247404432
```

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Set custom labels for confusion matrix
labels = ['No', 'Yes']
```

```
# Generate classification report and confusion matrix for Logistic Regression
logistic_report = classification_report(y_test, logistic_preds)
```

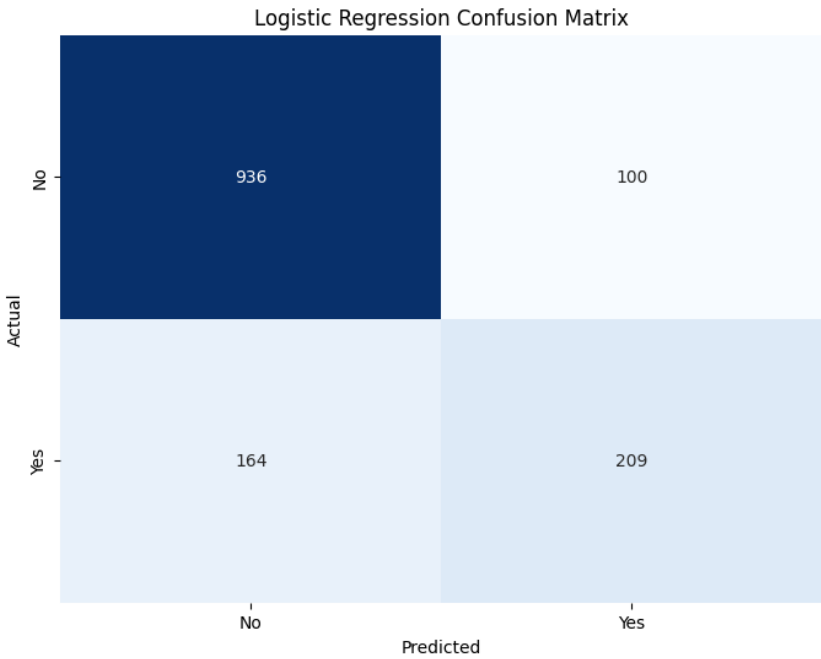
```
logistic_report = classification_report(y_test, logistic_preds)
logistic_conf_matrix = confusion_matrix(y_test, logistic_preds)

print("Logistic Regression Classification Report:")
print(logistic_report)

plt.figure(figsize=(8, 6))
sns.heatmap(logistic_conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False, xticklabels=labels, yticklabels=labels)
plt.title('Logistic Regression Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.85	0.90	0.88	1036
1	0.68	0.56	0.61	373
accuracy			0.81	1409
macro avg	0.76	0.73	0.74	1409
weighted avg	0.80	0.81	0.81	1409



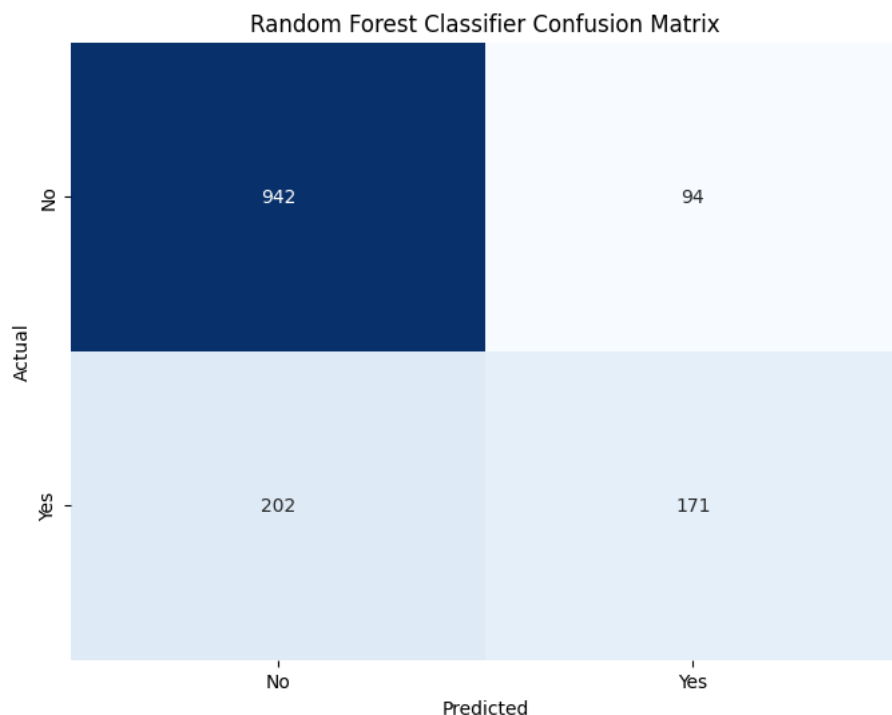
```
# Generate classification report and confusion matrix for Random Forest Classifier
rf_report = classification_report(y_test, rf_preds)
rf_conf_matrix = confusion_matrix(y_test, rf_preds)

print("\nRandom Forest Classifier Classification Report:")
print(rf_report)

plt.figure(figsize=(8, 6))
sns.heatmap(rf_conf_matrix, annot=True, cmap='Blues', fmt='g', cbar=False, xticklabels=labels, yticklabels=labels)
plt.title('Random Forest Classifier Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



Random Forest Classifier Classification Report:					
	precision	recall	f1-score	support	
0	0.82	0.91	0.86	1036	
1	0.65	0.46	0.54	373	
accuracy			0.79	1409	
macro avg	0.73	0.68	0.70	1409	
weighted avg	0.78	0.79	0.78	1409	



Model Tuning

Parameters used for tuning a logistic regression model:

1. `c`:

- It controls the regularization strength in logistic regression. Smaller values indicate stronger regularization, preventing overfitting by penalizing large parameter values.
- In the parameter grid, different values of `c` ranging from very small (0.001) to very large (100) are specified. This allows the grid search to explore a wide range of regularization strengths.

2. `penalty`:


- This parameter determines the type of regularization used in logistic regression.
- `'l1'` penalty refers to L1 regularization, which adds the absolute values of the coefficients to the loss function. It can lead to sparse solutions by pushing less informative features' coefficients to zero.
- `'l2'` penalty refers to L2 regularization, which adds the squared magnitudes of the coefficients to the loss function. It tends to shrink the coefficients towards zero without necessarily setting them to zero.
- By including both penalties in the parameter grid, the grid search will explore the effects of different types of regularization.

Overall, these parameters allow the grid search to systematically evaluate the logistic regression model's performance across various regularization strengths and types, helping to find the combination that optimizes the model's predictive performance.

```
from sklearn.model_selection import GridSearchCV
```

```
# Define parameter grid for logistic regression
logistic_param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2']
}
```

```
# Grid search with cross-validation for logistic regression
logistic_grid_search = GridSearchCV(LogisticRegression(), logistic_param_grid, cv=5, scoring='accuracy')
logistic_grid_search.fit(X_train, y_train)
```



Certainly! Here's an explanation of the parameters in the `rf_param_grid` dictionary used for tuning a Random Forest Classifier:

1. `n_estimators`:
 - This parameter determines the number of trees in the random forest. Each tree in the forest is built using a different random subset of the training data.
 - In the parameter grid, `[50, 100, 200]` are the candidate values for `n_estimators`. These values represent different numbers of trees to be included in the random forest.
2. `max_depth`:
 - The `max_depth` parameter controls the maximum depth of each decision tree in the random forest. A deeper tree can capture more complex relationships in the data, but it also increases the risk of overfitting.
 - In the parameter grid, `[None, 10, 20]` are the candidate values for `max_depth`. Using `None` means that there is no maximum depth limit, allowing the trees to grow until all leaves are pure or until they contain less than `min_samples_split` samples.
3. `min_samples_split`:
 - This parameter specifies the minimum number of samples required to split an internal node in a decision tree. It helps control the tree's complexity and prevents overfitting.
 - In the parameter grid, `[2, 5, 10]` are the candidate values for `min_samples_split`. These values represent different thresholds for splitting nodes based on the number of samples.
4. `min_samples_leaf`:
 - The `min_samples_leaf` parameter determines the minimum number of samples required to be at a leaf node. It helps prevent overfitting by controlling the minimum size of the leaves.
 - In the parameter grid, `[1, 2, 4]` are the candidate values for `min_samples_leaf`. These values represent different thresholds for the minimum number of samples required to form a leaf node.

```
# Define parameter grid for random forest classifier
rf_param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Grid search with cross-validation for random forest classifier
rf_grid_search = GridSearchCV(RandomForestClassifier(), rf_param_grid, cv=5, scoring='accuracy')
rf_grid_search.fit(X_train, y_train)
```



```
# Get best parameters and best scores
print("Best parameters for Logistic Regression:", logistic_grid_search.best_params_)
print("Best score for Logistic Regression:", logistic_grid_search.best_score_)

print("\nBest parameters for Random Forest Classifier:", rf_grid_search.best_params_)
print("Best score for Random Forest Classifier:", rf_grid_search.best_score_)

Best parameters for Logistic Regression: {'C': 1, 'penalty': 'l2'}
Best score for Logistic Regression: 0.8026256853811106

Best parameters for Random Forest Classifier: {'max_depth': 10, 'min_samples_leaf': 2, 'min_samples_split': 10, 'n_estimators': 200}
Best score for Random Forest Classifier: 0.8033355345381645
```

```
# Evaluate tuned models
tuned_logistic_model = logistic_grid_search.best_estimator_
tuned_rf_model = rf_grid_search.best_estimator_

tuned_logistic_preds = tuned_logistic_model.predict(X_test)
tuned_rf_preds = tuned_rf_model.predict(X_test)

# Compare performance before and after tuning
print("Logistic Regression Performance (Before Tuning):")
print(classification_report(y_test, logistic_preds))

print("\nLogistic Regression Performance (After Tuning):")
print(classification_report(y_test, tuned_logistic_preds))

print("\nRandom Forest Classifier Performance (Before Tuning):")
print(classification_report(y_test, rf_preds))

print("\nRandom Forest Classifier Performance (After Tuning):")
print(classification_report(y_test, tuned_rf_preds))
```

```
Logistic Regression Performance (Before Tuning):
```

	precision	recall	f1-score	support
0	0.85	0.90	0.88	1036
1	0.68	0.56	0.61	373
accuracy			0.81	1409
macro avg	0.76	0.73	0.74	1409
weighted avg	0.80	0.81	0.81	1409

```
Logistic Regression Performance (After Tuning):
```

	precision	recall	f1-score	support
0	0.85	0.90	0.88	1036
1	0.68	0.56	0.61	373
accuracy			0.81	1409
macro avg	0.76	0.73	0.74	1409
weighted avg	0.80	0.81	0.81	1409

```
Random Forest Classifier Performance (Before Tuning):
```

	precision	recall	f1-score	support
0	0.82	0.91	0.86	1036
1	0.65	0.46	0.54	373
accuracy			0.79	1409
macro avg	0.73	0.68	0.70	1409
weighted avg	0.78	0.79	0.78	1409

```
Random Forest Classifier Performance (After Tuning):
```

	precision	recall	f1-score	support
0	0.84	0.91	0.87	1036
1	0.68	0.51	0.58	373
accuracy			0.81	1409
macro avg	0.76	0.71	0.73	1409
weighted avg	0.79	0.81	0.80	1409

Interpretation and Conclusion

```
# 1. Interpretation of Model Results
# For Logistic Regression:
logistic_coefficients = logistic_model.coef_[0]
feature_names = X.columns

logistic_feature_importance = pd.DataFrame({'Feature': feature_names, 'Coefficient': logistic_coefficients})
logistic_feature_importance = logistic_feature_importance.sort_values(by='Coefficient', ascending=False)

print("Logistic Regression Coefficients:")
print(logistic_feature_importance)

# For Random Forest Classifier:
rf_feature_importance = pd.DataFrame({'Feature': feature_names, 'Importance': rf_model.feature_importances_})
rf_feature_importance = rf_feature_importance.sort_values(by='Importance', ascending=False)

print("\nRandom Forest Classifier Feature Importances:")
print(rf_feature_importance)
```

38	tenureGroup_Mid-term	-0.132405
3	Dependents	-0.137057
13	MonthlyCharges	-0.190163
5	PhoneService	-0.276744
26	MultipleLines_No	-0.300611
9	TechSupport	-0.325419
29	InternetService_DSL	-0.332883
31	InternetService_No	-0.337472
37	TenureGroup_Intermediate	-0.353014
6	OnlineSecurity	-0.373056
21	Contract_Two year	-0.850881
4	tenure	-3.420536

Random Forest Classifier Feature Importances:

	Feature	Importance
14	TotalCharges	0.153012
13	MonthlyCharges	0.142501
4	tenure	0.121467
19	Contract_Month-to-month	0.049192
6	OnlineSecurity	0.036864
9	TechSupport	0.035325
16	NumServices	0.031878
36	TenureGroup_New	0.030064
0	gender	0.026892
12	PaperlessBilling	0.024989
30	InternetService_Fiber optic	0.024271
7	OnlineBackup	0.022645
1	SeniorCitizen	0.020404
8	DeviceProtection	0.017300
2	Partner	0.017131
3	Dependents	0.016304
11	StreamingMovies	0.016110
15	FamilyStatus	0.015867
10	StreamingTV	0.015132
24	PaymentMethod_Electronic check	0.014850
21	Contract_Two year	0.014599
34	PaymentMethod_Electronic check	0.013669
26	MultipleLines_No	0.013007
28	MultipleLines_Yes	0.012994
29	InternetService_DSL	0.011142
40	TenureGroup_Very Long-term	0.009736
20	Contract_One year	0.008098
37	TenureGroup_Intermediate	0.007332
38	TenureGroup_Mid-term	0.007155
32	PaymentMethod_Bank transfer (automatic)	0.007073